

Homework 5 - COT5405

Kobee Raveendran

December 2021

1. **Suppose you have an undirected graph with only positive integer edge weights. How can you use BFS to find the length of the shortest path from a source vertex to every other vertex in $O((E + V)D)$ time, where D is the maximum weight on an edge.**

Since BFS is usually only useful for unweighted graphs (or graphs that have uniform edge weights), we can't apply it outright here. What we could instead do is "expand" all the weights. Since all edge weights here are positive integers, I'll say the "default" weight is 1. Thus, all edges with weight not equal to 1 will have to be expanded into a series of edges with weight 1. So, for example, an edge with weight 7 would be split up into 7 edges each with weight 1 (also with dummy vertices between these edges). Once we've done this for all the original edges, we can then run BFS as normal to find the shortest path lengths.

The only extra work we incur here is from splitting up all the edges with weights not equal to 1. This is a process bounded by the weight of the maximum edge in the graph, which we know here is D . Thus, for all $|E|$ edges, we incur at most D work per edge, and since inserting an edge means also inserting a vertex, for all vertices we also incur this D work per vertex. So, our final runtime is $O(|E|D + |V|D) = O((|E| + |V|)D)$.

2. **Consider a directed weighted graph with non-negative edge weights. For some cases, it is necessary to compute the length of the shortest path from every vertex v to a target vertex t . Describe how you can compute this in the same time complexity as in Dijkstra's algorithm.**

This problem seems like an inversion of the goal of Dijkstra's algorithm. To solve it, we can set the target vertex t in this problem to be the "source" vertex in Dijkstra's algorithm. Then, we just run through Dijkstra's algorithm to find the shortest paths from the "source" (our target t) to all other vertices. Since we don't need to maintain the actual paths, we do not need to reverse them (we can just maintain the length of the paths as we progress through). Since all we've done is run Dijkstra but swap the "source" and "target" vertices, the runtime is exactly the same.

3. **Consider a directed graph where each edge has some existence probability, i.e., say an edge e exists with probability $Pr(e)$, where $0 \leq Pr(e) \leq 1$. If a path from u to v has the edges e_1, e_2, \dots, e_k , then the probability that the path exists is given by $Pr(e_1) \cdot Pr(e_2) \cdot \dots \cdot Pr(e_k)$. Given a source vertex s and a target vertex t , describe how you can use Dijkstra's algorithm to find a path from s to t that has the maximum probability of existing.**

To obtain the maximum probability of a path existing, we should first try to maximize the probabilities of the constituent edges along the path, which will maximize their product. So, we should essentially be flipping another goal of Dijkstra's algorithm (which originally aims to find the shortest paths, but here we want the "longest" path in that the path's edge weights should be maximized rather than minimized). However, I don't believe Dijkstra's can be used to find longest paths outright, but we can try converting our problem from edge maximization to edge minimization. To do that, I'll use a trick from optimization commonly seen in machine learning loss function optimization, which is to seek to minimize the negative log likelihood instead of maximizing the regular likelihood. All we have to do is, for all edges in the graph, set the edge weight w_{edge} to $w_{edge} = -\log w_{edge}$.

We've now converted our maximization problem into a minimization one. If you aren't familiar with the negative log likelihood trick, to convince yourself that this works, see the graph of $-\log x$ in figure 1 below. Note how, in the range $[0, 1]$ (where our original edge weight probabilities lie), $-\log(x)$ is decreasing and non-negative as x increases. For example, $-\log(0.6) \approx 0.221$ while $-\log(0.8) \approx 0.097$; in other words, the higher probability has been

converted into a smaller value. Since we are also dealing with logs, we can replace the multiplication in the original problem with addition (so, the existence probability of a path is now the sum of its edge weights rather than the product).

Anyway, now that we have a minimization problem, we can just run Dijkstra's as normal on our updated graph and find the "shortest" path (the smallest constituent edge weights along the path). This path is the path that would yield the highest probability of existing in our original graph.

To do all this, we just need $O(|E|)$ time to first update each edge weight w to $-\log w$. Then we run Dijkstra's as normal, which of course doesn't change its runtime.




Figure 1: Graph of $f(x) = -\log x$. Note that our original edge probabilities (the values on the x -axis) are now mapped to a y -value on the curve. Also note that these y -values decrease from $+\infty$ to 0 as the x -values increase from 0 to 1.

4. **Given a directed acyclic graph, find the length of the longest path in the graph in $O(n+m)$ time. Note that we are NOT looking for the longest path starting at a particular vertex, rather the overall longest path that can start at any arbitrary vertex.**

NOTE: Here I'm assuming the graph is unweighted. If it is actually weighted, then we could possibly try another approach like negating the edge weights, then finding the shortest paths in this new graph.

In a DAG, we can find the longest path (and its length) starting at a vertex v using a DFS from that vertex. However, since we want the overall longest path, we'd have to run a DFS for all vertices, which would give us a quadratic runtime ($O(n^2)$, where $n = |V| + |E|$).

However, we can notice that these longest paths build upon each other and have an optimal substructure property, allowing us to use DP. For example, to get the length of the longest path from a node at one end of a graph to the other end, we need only add the length from that node to a central node and the length from that central node to a node at the opposite end of the graph.

In our DP matrix, the length of the longest path from all vertices will be initialized to 0. Then, iterating over the set of vertices, we assign $\text{dp}[v]$ to be the max over all of its neighbors lengths (plus 1) and its own value ($\text{dp}[v]$).

In other words, $\text{dp}[v] = \max(\max(\text{dp}[n_i] \text{ for } n_i \text{ in neighbors}) + 1, \text{dp}[v])$. Once we've run through this for all vertices, we need only scan through dp one final time and take the maximum value, which is our final answer. We are not doing any repeated work, and so in our loop over all $|V|$ vertices, we will just encounter all $|V|$ vertices and $|E|$ edges, giving us a total runtime of $O(|V| + |E|)$, or $O(n + m)$ as per the problem statement (using linear extra space for the DP table).